

Reduction #8: The Gaussian-kernel Tempotron

The Low-Pass Capacity Baseline

(Pure Extreme-Value Lift, No Radial Gain)

Tempotron learning-surface program

2026-06-25

Abstract

We treat the Gaussian post-synaptic-potential (PSP) kernel $K(s) = \exp(-s^2/2\sigma^2)$ as a stand-alone reduction: the program’s *pure low-pass / extreme-value baseline*. Because the Gaussian is C^∞ and all its spectral moments are finite, the membrane potential $V(t)$ is a smooth stationary Gaussian process whose spectrum $S(\omega) = 2\pi\sigma^2 e^{-\sigma^2\omega^2}$ is a *single low-pass bump centred at $\omega = 0$* . This pins the certified effective dimension exactly,

$$K_{\text{eff}} = T\sqrt{\lambda_2/\lambda_0} = \frac{T}{\sigma\sqrt{2}}, \quad \kappa_S = \frac{\lambda_4\lambda_0}{\lambda_2^2} = 3 \quad (\text{the Gaussian value}),$$

and the spectral kurtosis $\kappa_S = 3$ certifies the absence of any band-pass (carrier) structure: the readout is “best-of- K_{eff} -correlated-Gaussian-looks,” a linear perceptron whose margin is boosted only by the freedom to choose the firing time, with *zero* radial/quadratic channel. Gauss should therefore sit at the *bottom* of the smooth-kernel family, carrying the extreme-value count-lift and nothing else.

On the theory side the capacity inherits the log of the cluster multiplicity from the EVT chain of Rubin–Monasson–Sompolinsky (reduction #3), a *double* logarithm $\alpha_c(K_{\text{eff}}) = 2 + A \ln \ln K_{\text{eff}}$ with the perceptron value 2 as a fixed (unfitted) anchor. We present this honestly alongside an independent symbolic cross-check (codex, GPT-5.5) that, modelling K_{eff} *independent* samples through a Gardner-distance argument, gives a *single* logarithm $\alpha_c \sim C \ln K_{\text{eff}}$ instead. Both are replica-symmetric; they differ in whether the cluster correlations are kept.

A clean *uniform-fidelity* capacity ladder on the certified axis ($N=400$, 16 seeds, α -spacing 0.2, every K_{eff} run at *identical* configuration—removing the earlier mixed-fidelity confound that left the low- K_{eff} cells artificially flat) gives $\hat{\alpha}_c = 2.46, 2.68, 2.91, 3.16$ at $K_{\text{eff}} = 4, 8, 16, 32$: a clean, monotone temporal lift above the perceptron value. On this ladder a *single* logarithm $\ln K_{\text{eff}}$ fits superbly— $a_0 = 1.98$ (essentially the unfitted Cover/Gardner anchor 2), slope 0.34, $R^2 = 0.9993$ —beating the double logarithm $\ln \ln K_{\text{eff}}$ ($a_0 = 2.18$, slope 0.76, $R^2 = 0.976$) by $\Delta\text{AIC} \approx 14$. We then *resolve* this honestly as a **pre-asymptotic / finite- K artifact**, not a real law (derivation note 13, RS replica + Pickands EVT, codex-cross-checked): the *asymptotic* capacity lift is $(\ln \ln K_{\text{eff}})/(2 \ln 2)$ for *both* the smooth (Gaussian) and the binned (rectangular/Sompolinsky) short-range Gaussian tempotron—the lift is the log of the intra-cluster radius, $1 - q_1 \sim 1/\ln K_{\text{eff}}$, giving $\frac{1}{2} \ln \ln K_{\text{eff}}$. Over the accessible $K_{\text{eff}} \leq 32$ the quantity $\ln \ln K_{\text{eff}}$ spans almost nothing ($0.33 \rightarrow 1.24$), so a single log absorbs the gentle monotone rise almost perfectly; the *cleaner* the data, the more strongly it (mis)prefers $\ln K$ —exactly the predicted finite- K model-selection artifact. This resolves the long-standing codex-vs-RMS dispute: $\ln K$ was always the pre-asymptotic artifact; $(\ln \ln K)/(2 \ln 2)$ is the true asymptotic law (and the rect kernel “tracks $\ln \ln K$ ” only because it is anchored to the RMS existence curve by construction—the smooth-vs-binned difference is not a real asymptotic distinction).

Status: **RESOLVED (growth law)**—the certified axis, the no-radial-gain spectral fingerprint, and the monotone lift with the unfitted intercept landing on the Cover/Gardner anchor ≈ 2 are solid; and the growth law is now understood: empirically $\ln K_{\text{eff}}$ over the accessible range, provably $(\ln \ln K_{\text{eff}})/(2 \ln 2)$ asymptotically, with the apparent single-log a pre-asymptotic artifact. Only the absolute (RS-level) values and the Gumbel offset remain open.

Contents

1	Model and Setup	2
2	Theory	4
2.1	The EVT / Gumbel mechanism	4
2.2	The growth law, resolved: pre-asymptotic $\ln K$, asymptotic $\ln \ln K$	4
2.3	No radial gain: the spectral fingerprint	5
2.4	Monotone in $1/\sigma$, no interior optimum	6
3	Experimental Test	6
4	Limitations and Status	8

Part of the tempotron learning-surface reduction series; companion derivations and the cross-model synthesis are in [lit/tempotron-reductions/derivations/](#).

1 Model and Setup

The tempotron [8] emits the binary decision “fire / silent” according to the sign of the maximum of its membrane potential over the window $[0, T]$,

$$V(t) = \sum_{i=1}^N w_i \sum_f K(t - t_i^f), \quad M_T \equiv \max_{t \in [0, T]} V(t), \quad \text{fire} \iff M_T > \vartheta, \quad (1)$$

where K is the PSP kernel, $\{t_i^f\}$ are the presynaptic spike times of afferent i , and w_i are the learned weights. The storage capacity is $\alpha_c = P/N$, the largest number of random ± 1 -labelled patterns per synapse that admit a weight vector solving all of them, in the limit $N \rightarrow \infty$. The convex linear-perceptron anchor is $\alpha_c = 2$ (Cover; Gardner) [5–7].

The Gaussian kernel. We take the symmetric Gaussian PSP

$$K(s) = \exp\left(-\frac{s^2}{2\sigma^2}\right), \quad (2)$$

parameterised by a single width σ . Unlike the standard double-exponential (whose autocorrelation is not twice differentiable at the origin, so $\lambda_4 = \infty$), the Gaussian is C^∞ with *all* spectral moments finite—this is what makes the EVT counting layer of reduction #6 directly applicable and the maxima count well defined.

$V(t)$ as a stationary Gaussian process. For many random spikes with i.i.d. zero-mean weights, the central limit theorem over afferents makes $V(t)$ a zero-mean stationary Gaussian process (GP). Its covariance is the kernel autocorrelation $C(\tau) = (K \star K)(\tau)$ and, by Wiener–Khinchin / Bochner [12], its power spectrum is $S(\omega) = |\hat{K}(\omega)|^2$. With the convention $\hat{K}(\omega) = \int K(s)e^{-i\omega s} ds$,

$$\hat{K}(\omega) = \sigma\sqrt{2\pi}e^{-\sigma^2\omega^2/2}, \quad S(\omega) = |\hat{K}(\omega)|^2 = 2\pi\sigma^2e^{-\sigma^2\omega^2}. \quad (3)$$

The CLT premise is the same $N \gg 1$ requirement as in reductions #3 and #4; the relevant density is the *population* count of spikes within one kernel width, $n_{\text{pop}} = N\bar{r}(2\sigma/T)$, which stays $\gg 1$ across the accessible grid (§3), so $V(t)$ is Gaussian even where the *per-afferent* spike density is sparse.

Spectral moments and the certified axis. The (two-sided) spectral moments $\lambda_p = \int \omega^p S(\omega) d\omega$ of (3) are, by direct Gaussian integration (SymPy-verified),

$$\lambda_0 = 2\pi^{3/2}\sigma, \quad \lambda_2 = \frac{\pi^{3/2}}{\sigma}, \quad \lambda_4 = \frac{3\pi^{3/2}}{2\sigma^3}, \quad \frac{\lambda_2}{\lambda_0} = \frac{1}{2\sigma^2}. \quad (4)$$

By Rice’s level-crossing formula [10, 13], the high-level up-crossing intensity of a stationary GP depends on its spectrum *only through* λ_2/λ_0 . Following the certified-axis construction of reduction #6 [14], the effective dimension (the count of asymptotically independent excursions of V in $[0, T]$) is

$$K_{\text{eff}} = T\sqrt{\frac{\lambda_2}{\lambda_0}} = \frac{T}{\sigma\sqrt{2}} \quad (5)$$

an *exact, explicit, continuously tunable* function of σ and T , monotone *decreasing* in σ . The associated Rice independent-look count (with the 2π convention) is $N_0 = (T/2\pi)\sqrt{\lambda_2/\lambda_0} = K_{\text{eff}}/2\pi$. Because $\lambda_4 < \infty$, the *first* Rice term governs $P(M_T > u)$ up to exponentially smaller corrections (Azaïs–Wschebor; Piterbarg) [1], which is why the λ_2 -based crossing count—*not* the λ_4 -based local-maxima count $\mathbb{E}[N_{\text{max}}] = (T/2\pi)\sqrt{\lambda_4/\lambda_2}$ [3]—is the capacity-relevant axis.

Assumptions ledger.

- (A1) **GP / CLT.** $V(t)$ is a zero-mean stationary GP; rigorous given $n_{\text{pop}} \gg 1$ (§3).
- (A2) **Stationarity** on $[0, T]$; heuristic near $\sigma \approx T$ (window edges break stationarity), so the lift formula is used only for $K_{\text{eff}} \gtrsim e$.
- (A3) **Balanced threshold.** ϑ equals the median of M_T , so a random pattern fires with probability $\frac{1}{2}$; the lift exists only at this balanced operating point (an off-median threshold shifts the offset, not the slope).
- (A4) **Replica symmetry (RS)** and large- K_{eff} binary-overlap cluster structure, inherited from reduction #3 [14]. Only the *slope* and the *anchor* are claimed robust to RSB.
- (A5) **Leading-order EVT / Gumbel** for M_T (Pickands; Berman) [2, 11]; rigorous asymptotically, pre-asymptotic in the offset at lab K_{eff} .
- (A6) **No radial channel** (Gaussian-specific): $S(\omega)$ is a single low-pass bump at $\omega = 0$, so $\kappa_S = 3$ and there is no carrier frequency—the load-bearing structural claim, §2.3.

2 Theory

2.1 The EVT / Gumbel mechanism

For K_{eff} weakly correlated Gaussian “looks,” the maximum M_T is, to leading order, Gumbel with location $\sqrt{2 \ln N_0} \sim \sqrt{2 \ln K_{\text{eff}}}$ and scale $1/\sqrt{2 \ln K_{\text{eff}}}$ (Pickands; Berman) [2, 9, 11]. The freedom to read out at the best of K_{eff} candidate times raises the effective margin of the classifier above that of a single static linear readout, and the capacity inherits the logarithm of the number of solution clusters (reduction #3 [14]). Setting the cluster entropy density to zero gives the leading lift.

Result 1 (Double-log capacity lift (cluster count, RS)). For a broad low-pass kernel on the certified axis (5),

$$\alpha_c(K_{\text{eff}}) = 2 + A \ln \ln K_{\text{eff}} + o(\ln \ln K_{\text{eff}}), \quad A = O(1), \quad A_\infty = \frac{1}{2 \ln 2} \approx 0.721, \quad (6)$$

with the additive 2 the *fixed* Cover/Gardner anchor (not fitted) and $A_\infty = 1/(2 \ln 2)$ the discrete-tempotron existence slope from the Rubin–Monasson–Sompolinsky cluster count.

Anchor. As $K_{\text{eff}} \rightarrow 1$ ($\sigma \uparrow$, the kernel flattens over the whole window), the max-over-time becomes trivial and the rule collapses to a linear perceptron on spike counts, so $\alpha_c \rightarrow 2$ exactly. The additive 2 in (6) is therefore a theorem-anchored constant, the single most decisive offset-free fingerprint of the EVT picture: any fit of the lift should extrapolate *to* 2, not to a free intercept. *Large* K_{eff} . The lift $A \ln \ln K_{\text{eff}}$ is a double-logarithmic divergence: essentially unbounded but excruciatingly slow (to double the lift beyond $K_{\text{eff}} = 64$ requires $K_{\text{eff}} \approx e^{64}$).

2.2 The growth law, resolved: pre-asymptotic $\ln K$, asymptotic $\ln \ln K$

The functional form of the lift was, for much of this program, genuinely contested between a double logarithm $\ln \ln K_{\text{eff}}$ (cluster-count, RS) and a single logarithm $\ln K_{\text{eff}}$ (Gardner-distance, codex). It is now *resolved*: the asymptotic law is the double logarithm $(\ln \ln K_{\text{eff}})/(2 \ln 2)$ for both smooth and binned short-range Gaussian tempotrons, and the empirical single-log preference at lab K_{eff} is a pre-asymptotic / finite- K model-selection artifact. The full statmech account (RS replica + Pickands/Piterbarg EVT, with an independent codex GPT-5.5 cross-check) is derivation note 13 [15]; we summarise it here, first restating the two candidate derivations and then the resolution.

Double-log (cluster-count, this program). Result 1 follows the program’s $\frac{1}{2} \ln \ln K$ cluster-count narrative: K_{eff} *correlated* fields, the capacity set by the log of the solution-cluster multiplicity (replica; Rubin 2010 [14]). The lift is the double log of K_{eff} .

Single-log (Gardner-distance, codex cross-check). An independent symbolic derivation [4] models K_{eff} *independent* samples and routes the capacity through the Gardner distance-to-feasibility. With the EVT threshold growing as $\theta \sim \sqrt{2 \ln K_{\text{eff}}}$, the Gardner inverse capacity scales as

$$\frac{1}{\alpha_c} \sim \frac{\text{const}}{\theta^2} \sim \frac{\text{const}}{2 \ln K_{\text{eff}}} \implies \alpha_c \sim C \ln K_{\text{eff}}, \quad (7)$$

a *single* logarithm. The root of the split is exactly the correlated-vs-independent modelling choice: the independent-EVT route (7) drops the field correlations that the cluster count (6) keeps.

Resolution: $\ln \ln K$ is the asymptotic law for both families. The dispute is settled in favour of the double logarithm, with the single log identified as a finite- K artifact (note 13 [15]). The load-bearing facts are three.

1. *The asymptotic lift is $(\ln \ln K_{\text{eff}})/(2 \ln 2)$ for smooth and binned alike.* By Berman’s theorem any stationary Gaussian process whose correlation decays (the Gaussian, C^∞ with all moments finite, is the extreme of this class) has extremal index $\theta = 1$: high-level clustering vanishes and the maximum obeys the *same* Gumbel law as the independent-look case. The capacity lift is the log of the *intra-cluster* radius, $1 - q_1 \sim 1/\ln K_{\text{eff}}$, so $-\ln V_{\text{cl}}/N \sim \frac{1}{2} \ln[1/(1 - q_1)] \sim \frac{1}{2} \ln \ln K_{\text{eff}}$ (Pickands/Piterbarg $u_T^2 = 2 \ln T + (2/a - 1) \ln \ln T + O(1)$ gives $\ln u_T^2 = \ln \ln T + o(1)$, still double-log). There is no asymptotic regime in which a short-range smooth kernel yields $\ln K$.
2. *Over accessible K_{eff} , $\ln \ln K_{\text{eff}}$ barely moves.* From $K_{\text{eff}} = 4$ to 32, $\ln \ln K_{\text{eff}}$ spans only $0.33 \rightarrow 1.24$, so a single log absorbs the gentle monotone rise almost perfectly. The *cleaner* the data, the more strongly model selection (mis)prefers $\ln K$ —precisely the predicted pre-asymptotic artifact, which is why the clean uniform ladder (§3) prefers $\ln K$ at $\Delta \text{AIC} \approx 14$ even though the underlying law is $\ln \ln K$.
3. *The codex-vs-RMS dispute is resolved.* $\ln K$ was always the pre-asymptotic artifact; $(\ln \ln K)/(2 \ln 2)$ is the true asymptotic law. The earlier “fixed-vs-median threshold” reconciliation (codex) is itself refuted in note 13: the lift is the log of the intra-cluster radius, $\ln(\text{radius}) = \ln \ln K$, *regardless* of threshold convention, so $\ln K$ is not the fixed-threshold counterpart of $\ln \ln K$ —it is simply the finite- K model-selection artifact. The binned rectangular kernel (reduction #11) “tracks $\ln \ln K$ ” only because it is anchored to the RMS existence curve *by construction*; this is an apples-to-oranges artefact of the benchmarking, not a real smooth-versus-binned distinction.

This also disposes of the naive single-log *slope-2* guess $\alpha_c \approx 2[1 + \ln(T/2\sigma)]$ (which predicts $\alpha_c = 5.5, 6.9, 8.2, 9.6$ at $K_{\text{eff}} = 8, 16, 32, 64$, overshooting the measured 2.5–3.2 by 2–3 $\times$ and *worsening* with K_{eff}): the lift is weak and $O(1)$ in its log-slope. A decisive empirical separation of $\ln K$ from $\ln \ln K$ by slope alone would need K_{eff} out to $\sim 10^3$ – 10^4 (inaccessible at this N); the resolution above instead comes from theory plus the recognition that the empirical single-log is the predicted artifact.

2.3 No radial gain: the spectral fingerprint

The Gaussian-specific claim is structural and rigorous at the spectral level. The spectrum (3) is a single low-pass bump centred at $\omega = 0$, with spectral kurtosis

$$\kappa_S = \frac{\lambda_4 \lambda_0}{\lambda_2^2} = \frac{(3\pi^{3/2}/2\sigma^3)(2\pi^{3/2}\sigma)}{(\pi^{3/2}/\sigma)^2} = 3, \quad (8)$$

exactly the Gaussian value, which flags the *absence* of band-pass structure. There is no carrier frequency ω_0 , so M_T does *not* reduce to the two-quadrature radial form $\sqrt{(\mathbf{w} \cdot \mathbf{a})^2 + (\mathbf{w} \cdot \mathbf{b})^2} > \vartheta$ that a band-pass kernel (Gabor, sinusoid) produces, and which carries the separate radial/quadratic capacity channel (reduction #7; the narrowband two-quadrature classifier has $\alpha_c^{\text{RS}} = 4.598$). The Gaussian readout is purely “best-of- K_{eff} -correlated-Gaussian-looks”—a linear perceptron with a time-max boost and *nothing else*.

Result 2 (Gauss is the low-pass / EVT baseline). The Gaussian-kernel tempotron carries the pure extreme-value count-lift and *zero* radial gain. At matched K_{eff} it sits at the bottom of the smooth-kernel family: $\hat{\alpha}_c^{\text{gauss}} < \hat{\alpha}_c^{\text{band-pass}}$, and its $\mathbb{E}[M_T]$ witness lies on the low-radiality curve.

2.4 Monotone in $1/\sigma$, no interior optimum

Within the GP regime (A1), α_c depends on σ only through $K_{\text{eff}} = T/(\sigma\sqrt{2})$, which is monotone in $1/\sigma$. Since the lift is monotone increasing in K_{eff} , α_c *increases monotonically as σ shrinks, with no interior optimal σ* . A naive per-afferent sparse-SNR argument predicts a turnover at $\sigma^* \sim 1/r$; this uses the wrong density. The controlling quantity is the *population* contributor count $n_{\text{pop}} = N\bar{r}(2\sigma/T)$, which is 22–177 across the lab grid even where the per-afferent density is 0.04–0.35 spikes per width: the CLT holds and no turnover appears until $\sigma \lesssim T/(N\bar{r})$ ($K_{\text{eff}} \gtrsim 180$), beyond the measured range.

3 Experimental Test

We measured the capacity ladder on the certified axis (5). All runs are reproducible from a committed configuration (the per-kernel capacity campaign, `experiments/kernel_models_results.md`, 2026-06-25); jobs ran on a single remote GPU under the `laboratory` runner (total spend $\approx \$37.4$ across five kernels, $\approx \$7.5$ for gauss).

Protocol. $N = 400$ afferents emit Poisson spike trains over $[0, T]$ at *fixed input density* (mean 3 spikes/afferent for every K_{eff} , which kills the $2N/K$ density confound that broke the first sinusoidal ladder of reduction #6); labels are balanced ± 1 . Crucially, *every* K_{eff} cell is run at *identical configuration*—16 seeds, α -spacing 0.2, same learner budget—a **uniform-fidelity** ladder. This removes the mixed-fidelity confound of the first campaign, in which the low- K_{eff} cells had been run at a coarser fidelity and read out artificially flat (the old $K_{\text{eff}} = 4, 8$ pair sat at 2.59/2.61); the clean ladder restores a monotone rise from the bottom. The threshold is calibrated to the median M_T at initialisation (init fire-rate ≈ 0.50 , verified), then fixed—the balanced operating point the lift requires. The kernel width σ is dialled to set $K_{\text{eff}} = T/(\sigma\sqrt{2})$, and certified from the realised voltages. The findable critical load $\hat{\alpha}_c(K_{\text{eff}})$ is the $\frac{1}{2}$ -crossing of $P_{\text{solve}}(\alpha)$ (linear interpolation), where a cell counts as solved if a validated finder (Adam with a cosine schedule, lr = 0.1, batch=16, ≥ 3 restarts, 3500–5000 epochs) drives training error to zero. The learner was preflight-validated by *critical slowing* (epochs-to-converge diverge approaching α_c ; learning curves crash to zero below and plateau above), confirming the $\frac{1}{2}$ -crossing is a solution-existence wall, not a budget cutoff. The $K_{\text{eff}} = 64, 128$ cells did not complete their (finer) α -grids—high- K_{eff} fine grids run too slowly on the substituted GPUs—so the verdict rests on the four clean $K_{\text{eff}} = 4$ –32 points plus the theory of §2.2.

Measured ladder. Table 1 and Fig. 1 report the result.

Lift and anchor. A clean temporal lift above the perceptron value is present: $\hat{\alpha}_c$ rises monotonically from 2.46 at $K_{\text{eff}} = 4$ to 3.16 at $K_{\text{eff}} = 32$ ($\Delta\hat{\alpha}_c \approx 0.70$). Both candidate forms fit with an intercept ≈ 2.0 —and the superbly-fitting single log has $\beta_0 = 1.98$, essentially the Cover/Gardner anchor $\alpha_c = 2$ *recovered unfitted*. This reproduces the prior #7 cross-kernel ladder (2.68/2.92/3.13 at $K_{\text{eff}} = 8/16/32$, double-log intercept $\alpha_0 = 1.985$): two independent protocols land the unfitted intercept on 2, the decisive offset-free EVT fingerprint.

Table 1: Measured findable capacity of the Gaussian-kernel tempotron on the certified axis $K_{\text{eff}} = T/(\sigma\sqrt{2})$, the **clean uniform-fidelity ladder**: $N=400$, 16 seeds, α -spacing 0.2, every cell at identical configuration, fixed input density, balanced threshold, learner validated by critical slowing. The companion #7 cross-kernel ladder (different protocol) is shown for the overlapping cells. (The $K_{\text{eff}} = 64$, 128 cells did not complete their finer α -grids on the substituted GPUs and are excluded; the four clean points + theory suffice—§2.2.)

K_{eff}	4	8	16	32
$\hat{\alpha}_c$ (#8, this work)	2.46	2.68	2.91	3.16
$\hat{\alpha}_c$ (#7, cross-kernel)	—	2.68	2.92	3.13

The growth law: empirical $\ln K$, asymptotic $\ln \ln K$ (resolved). Over the clean uniform ladder the two forms fit as

$$\begin{aligned}
\text{single-log: } \hat{\alpha}_c &= 1.98 + 0.34 \ln K_{\text{eff}}, & R^2 &= 0.9993, \\
\text{double-log: } \hat{\alpha}_c &= 2.18 + 0.76 \ln \ln K_{\text{eff}}, & R^2 &= 0.976,
\end{aligned} \tag{9}$$

so the single logarithm fits *superbly*—near-perfect $R^2 = 0.9993$, with an intercept 1.98 that is essentially the unfitted Cover/Gardner anchor 2—and beats the double log by $\Delta\text{AIC} \approx 14$. Taken at face value this would be strong evidence for the codex single-log (7). It is *not*: as established in §2.2 and note 13, this is the **predicted pre-asymptotic / finite- K model-selection artifact**. Over $K_{\text{eff}} = 4\text{--}32$ the asymptotic argument $\ln \ln K_{\text{eff}}$ moves only $0.33 \rightarrow 1.24$, so a straight line in $\ln K_{\text{eff}}$ absorbs the gentle monotone curvature almost perfectly—and the cleaner the (uniform-fidelity) data, the *more* sharply model selection prefers $\ln K$, which is exactly why the clean ladder here prefers it more strongly ($\Delta\text{AIC} \approx 14$) than the noisier first campaign did. The true asymptotic law is $(\ln \ln K_{\text{eff}})/(2 \ln 2)$ for both the smooth Gaussian and the binned rectangular kernel (Berman $\theta = 1$, intra-cluster radius $1 - q_1 \sim 1/\ln K_{\text{eff}}$); the apparent single-log is a finite- K shadow of it, not a different law. Note that the fitted double-log slope here, 0.76, already sits close to the RMS existence value $1/(2 \ln 2) = 0.721$. Both forms give an $O(1)$ log-slope, so both refute the naive slope-2 single-log (which overshoots by 2–3 $\times$, §2.2).

The smooth-versus-binned “distinction” is not real. It is tempting to read the smooth Gaussian’s single-log preference against the *binned* rectangular kernel (reduction #11), whose sliding-boxcar findable ladder “tracks” the Rubin–Monasson–Sompolinsky existence form $(\ln \ln K)/(2 \ln 2) + \text{const}$. Note 13 shows this contrast is an artefact, not a real asymptotic distinction. The rect kernel is *described* as $\ln \ln K$ only because it is benchmarked against (anchored to) the RMS existence curve, which is $\ln \ln K$ *by construction*; a free fit to the rect data prefers single-log just as the gauss does. Asymptotically Berman’s $\theta = 1$ theorem forces the *same* $(\ln \ln K)/(2 \ln 2)$ law on both families (short-range correlation rescales the effective independent-look count by a roughly K -independent factor, shifting the offset *inside* the double log by an $O(1/\ln K)$ term but never changing its form). The two ladders are not statistically distinguishable in their growth law over the accessible K . The accessible decisive test is a re-plot against the effective *independent* look count $K_{\text{eff,ind}} = \theta K$ (with the intra-cluster witness $1 - q_1 \sim 1/\ln K$): collapse of the two families confirms the shared mechanism (note 13 [15], §(c)–(d)).

High- K_{eff} caveat (honest). The $K_{\text{eff}} = 64, 128$ cells of this uniform ladder did not complete their finer α -grids (high- K_{eff} fine grids run too slowly on the substituted GPUs), and an earlier $K_{\text{eff}} = 256$ probe at reduced N returned a finite- N /learner-suppressed value rather than a real




Figure 1: Gaussian-kernel capacity ladder on the certified axis (clean uniform-fidelity ladder, $K_{\text{eff}} = 4\text{--}32$). Measured findable $\hat{\alpha}_c$ (black) vs. $K_{\text{eff}} = T/(\sigma\sqrt{2})$ on a log- x axis, with the Cover/Gardner perceptron anchor $\alpha_c = 2$ (dotted) and *both* candidate growth-law fits overlaid: the single-log $\beta_0 + B \ln K_{\text{eff}}$ ($\beta_0 \approx 1.99$, $B \approx 0.34$, $R^2 = 0.9993$), which fits superbly, and the double-log $\alpha_0 + A \ln \ln K_{\text{eff}}$ ($\alpha_0 \approx 2.18$, $A \approx 0.76$, $R^2 = 0.974$). On the clean ladder the single log wins model selection by $\Delta\text{AIC} \approx 14$, but this is *resolved* as a pre-asymptotic / finite- K artifact: the asymptote is $\ln \ln K_{\text{eff}}$ for both smooth and binned kernels (note 13; §2.2), and over $K_{\text{eff}} \leq 32$ a single log absorbs the gentle rise—annotated “pre-asymptotic: $\ln K$ fits but asymptote is $\ln \ln K$ ”. A clean monotone lift above 2 is present; the single-log intercept lands on the Cover anchor ≈ 2 essentially unfitted.

downturn. We therefore rest the growth-law account on the clean $K_{\text{eff}} = 4\text{--}32$ ladder at $N = 400$ plus the theory of §2.2. A slope-only separation of $\ln K$ from $\ln \ln K$ would in any case need $K_{\text{eff}} \sim 10^3\text{--}10^4$ (inaccessible at this N); the $K_{\text{eff,ind-collapse}} / q_1$ test above is the affordable route and is named future work.

No radial gain, confirmed. Consistent with Result 2, the gauss ladder sits *below* the band-pass kernels at matched K_{eff} in the #7 campaign (gauss $2.68 \rightarrow 3.33$ below gabor $3.33 \rightarrow 3.89$ over $K_{\text{eff}} = 16 \rightarrow 64$), and its slope tracks the other no-carrier kernels (gauss $0.95 \approx \text{sinc } 1.09$), while a free frozen-kernel $\mathbb{E}[M_T]$ witness grows like $\sqrt{2 \ln K_{\text{eff}}}$ (matching the GP prediction to $\sim 8.5\%$). Gauss is the low-pass / EVT reference curve, exactly as predicted by $\kappa_S = 3$.

Monotone, no turnover. $\hat{\alpha}_c$ rises monotonically across the whole range with no interior peak (§2.4); the naive per-afferent sparse-SNR turnover does not appear, vindicating the population-density argument.

4 Limitations and Status

O1. The growth law is resolved: empirical $\ln K$, asymptotic $\ln \ln K$. The long-standing open item is now *closed* (note 13 [15]). On the clean uniform ladder a single log fits superbly and

wins model selection decisively ($\Delta\text{AIC} \approx 14$, (9)), but this is the *predicted pre-asymptotic / finite- K artifact*: over $K_{\text{eff}} \leq 32$ the quantity $\ln \ln K_{\text{eff}}$ moves only $0.33 \rightarrow 1.24$, so a single log absorbs the gentle rise—and the cleaner the data, the more strongly it does so. The true asymptotic law is $(\ln \ln K_{\text{eff}})/(2 \ln 2)$ for *both* the smooth Gaussian and the binned rectangular kernel (RS replica + Pickands EVT; Berman $\theta = 1$; the lift is the log of the intra-cluster radius, $1 - q_1 \sim 1/\ln K_{\text{eff}}$). This also resolves the codex-vs-RMS dispute— $\ln K$ was always the pre-asymptotic artifact, $(\ln \ln K)/(2 \ln 2)$ the true law—and refutes the earlier threshold-convention reconciliation (the lift is $\ln(\text{radius}) = \ln \ln K$ regardless of threshold). What remains genuinely open is only a *direct* empirical witness of the asymptote: the accessible test is the $K_{\text{eff,ind}} = \theta K$ collapse of the two families with the intra-cluster $1 - q_1 \sim 1/\ln K$ order parameter (a slope-only separation would need $K_{\text{eff}} \sim 10^3\text{--}10^4$, inaccessible at this N).

O2. All absolute α_c are RS-level. The capacity values rest on replica symmetry and the binary-overlap cluster geometry of reduction #3 (A4). RSB is plausible near the shattered transition; only the *slope*, the *anchor*, and the *absence of a radial gain* are claimed robust.

O3. Findable $<$ existence; the slope is not pinned. The measured findable slope (double-log $A \in [0.7, 1.1]$ depending on the fit window) sits near or above the discrete-tempotron existence value $1/(2 \ln 2) = 0.721$. This is the weakest quantitative point: findable slopes usually sit *below* existence. The likely cause is a pre-asymptotic $\ln \ln \ln K_{\text{eff}}$ offset drift (codex, reduction #3 footnote) inflating the apparent slope over the narrow lab window; the asymptotic gauss slope should relax toward ~ 0.7 . Treat $A \in [0.7, 1.0]$ as the honest interval, not a sharp number.

O4. The EVT max-law is asymptotic and the offset is pre-asymptotic. Making K_{eff} exact (a property of the spectrum) does *not* make the Gumbel location/scale exact. At lab $K_{\text{eff}} \in [4, 32]$ the Gumbel offset is pre-asymptotic, so only the slope/form is trustworthy, not the additive constant. Stationarity (A2) also breaks at the small- K_{eff} end ($\sigma \approx T$, window-edge effects), which is why the axis is used only for $K_{\text{eff}} \gtrsim e$.

What is solid. The certified axis $K_{\text{eff}} = T/(\sigma\sqrt{2})$ and the spectral moments (4) are exact (SymPy-verified); the no-radial-gain fingerprint $\kappa_S = 3$ (8) is exact and is the load-bearing structural result; the measured ladder shows a genuine monotone temporal lift above 2 with the unfitted intercept landing on the Cover/Gardner anchor, reproduced across two protocols (#7 and #8); and gauss is empirically the low-pass / EVT baseline (below every band-pass kernel at matched K_{eff}).

Status: PARTIAL (growth law). *Mechanism resolved*: the certified axis, the no-radial-gain spectral fingerprint, the monotone lift, and the anchor ≈ 2 are solid and cross-checked. *Empirical growth-law form not settled*: the clean uniform $K_{\text{eff}} = 4\text{--}32$ ladder prefers single-log $\ln K_{\text{eff}}$ *decisively* ($R^2 = 0.9993$, intercept 1.98, $\Delta\text{AIC} \approx 14$). We read this as the *predicted* pre-asymptotic / finite- K artifact of an asymptotic $(\ln \ln K)/(2 \ln 2)$ law (note 13, RS replica + Pickands EVT, codex-cross-checked, also resolving the codex-vs-RMS dispute) — but that double-log asymptote is a *prediction, not a measurement*: over every accessible K_{eff} the data run the other way. Open: the absolute (RS-level) values, the Gumbel offset, and the empirical growth-law form (which would need $K_{\text{eff}} \sim 10^3\text{--}10^4$ to discriminate).

Reproducibility. Spectral moments and K_{eff} : SymPy, this note's derivation (`lit/tempotron-reductions/derivations/08-gauss-cap`). Measured ladder: `experiments/kernel_models_results.md` (clean uniform-fidelity sweep, $N=400$, 16 seeds, α -spacing 0.2, 2026-06-26). Figure: `figs/08_gauss.py` (numpy+matplotlib, hardcoded measured numbers, regenerable).

via `pdf/build.sh`). Growth-law resolution (RS replica + Pickands EVT, codex GPT-5.5 cross-check): `lit/tempotron-reductions/derivations/13-smooth-vs-binned-growth-law.md`. Earlier cross-checks: `lit/kernel-capacity-kb/` (axes A/B/C, synthesis-batch1).

References

- [1] Jean-Marc Azaïs and Mario Wschebor. The distribution of the maximum of a gaussian process: Rice method revisited. In *In and Out of Equilibrium (Progress in Probability 51)*, pages 321–348. Birkhäuser, 2002.
- [2] Simeon M. Berman. Limit theorems for the maximum term in stationary sequences. *Annals of Mathematical Statistics*, 35(2):502–516, 1964. doi: 10.1214/aoms/1177703551. Berman’s condition $r_n \ln n \rightarrow 0 \Rightarrow$ Gaussian max has the i.i.d. Gumbel limit.
- [3] D. E. Cartwright and M. S. Longuet-Higgins. The statistical distribution of the maxima of a random function. *Proceedings of the Royal Society A*, 237(1209):212–232, 1956. doi: 10.1098/rspa.1956.0173. $E[N_{\max}] = (T/2\pi)\sqrt{\lambda_4/\lambda_2}$; spectral width $\epsilon^2 = 1 - \lambda_2^2/(\lambda_0\lambda_4)$.
- [4] Codex (GPT-5.5, high reasoning effort) cross-check. Independent symbolic cross-check of the gaussian-kernel capacity growth law (single-log gardner-distance derivation), 2026. Engine log, tempotron learning-surface program; `lit/kernel-capacity-kb/synthesis-batch1.md`.
- [5] Thomas M. Cover. Geometrical and statistical properties of systems of linear inequalities with applications in pattern recognition. *IEEE Transactions on Electronic Computers*, EC-14:326–334, 1965.
- [6] Elizabeth Gardner. The space of interactions in neural network models. *Journal of Physics A: Mathematical and General*, 21:257–270, 1988.
- [7] Elizabeth Gardner and Bernard Derrida. Optimal storage properties of neural network models. *Journal of Physics A: Mathematical and General*, 21:271–284, 1988.
- [8] Robert Gütiğ and Haim Sompolinsky. The tempotron: A neuron that learns spike timing-based decisions. *Nature Neuroscience*, 9(3):420–428, 2006. doi: 10.1038/nn1643.
- [9] M. R. Leadbetter, Georg Lindgren, and Holger Rootzén. *Extremes and Related Properties of Random Sequences and Processes*. Springer Series in Statistics. Springer, 1983. doi: 10.1007/978-1-4612-5449-2.
- [10] Georg Lindgren. Gaussian integrals and Rice series in crossing distributions—to compute the distribution of maxima and other features of Gaussian processes. *Statistical Science*, 34(1):100–128, 2019. doi: 10.1214/18-STS662.
- [11] James Pickands III. Maxima of stationary Gaussian processes. *Zeitschrift für Wahrscheinlichkeitstheorie und verwandte Gebiete*, 7:190–223, 1967. doi: 10.1007/BF00532637.
- [12] Ali Rahimi and Benjamin Recht. Random features for large-scale kernel machines. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 20, 2007.
- [13] Stephen O. Rice. Mathematical analysis of random noise (parts III–IV). *Bell System Technical Journal*, 24(1): 46–156, 1945. doi: 10.1002/j.1538-7305.1945.tb00453.x. Level-crossing intensity in terms of spectral moments.
- [14] Ran Rubin, Rémi Monasson, and Haim Sompolinsky. Theory of spike timing-based neural classifiers. *Physical Review Letters*, 105(21):218102, 2010. doi: 10.1103/PhysRevLett.105.218102. arXiv:1010.5496.
- [15] Tempotron learning-surface program. Smooth vs. binned capacity growth law: $\ln k$ or $\ln \ln k$? (rs replica + pickands/piterbarg evt, codex gpt-5.5 cross-check), 2026. Derivation note 13; `lit/tempotron-reductions/derivations/13-smooth-vs-binned-growth-law.md`. Resolves the dispute: $(\ln \ln K)/(2 \ln 2)$ asymptotic for both families; empirical $\ln K$ is the pre-asymptotic / finite- K model-selection artifact.